

Requirements Document — Leak-Proof ML Trading Signals on Minute-Bar Crypto

0) Audience & Purpose

This document specifies exactly what to build, how to build it, and why—so an engineer or an AI assistant can implement a **deployable** trading signal generator for crypto (minute bars) that **avoids overfitting** and **generalizes**. It covers the end-to-end pipeline: data → features → labels → validation → models → calibration → ensembling → trade construction → risk → backtesting → deployment → monitoring → retraining. It leaves no room for ambiguity and encodes guardrails against common finance ML errors (leakage, non-stationarity, backtest mirages).

Hard constraints:

- **No data leakage** (strict causality at every step).
 - **All metrics after fees/slippage** only.
 - **Validation must be purged & embargoed, walk-forward, and nested for hyperparameters.**
 - **Signals must be calibrated probabilities and mapped to expected value (EV), with hysteresis and risk limits.**
-

1) Problem Context & Goals

1.1 Context

- Data: 1-minute OHLCV + trades for a single crypto asset (extendable to many). High noise, regime shifts, heavy tails, non-stationarity.
- Current use: classification to produce **Long / Short / Neutral** signals; probabilities used as confidence.

1.2 Primary Goal

Produce **clear, tradable entry/exit signals** with a **repeatable statistical edge** that persists out-of-sample/live, measured **after realistic costs**.

1.3 Secondary Goals

- Robustness across regimes; low sensitivity to hyperparameters.
 - Controlled turnover and drawdowns.
 - Transparent, auditable pipeline with reproducible research and deployment.
-

2) Definitions & Notation

- tt : bar index; all features at tt use information **available at or before** bar close tt .
- Price $P_{t,t}$: use **close** unless specified.
- Log return: $r_t = \ln P_{t,t} - \ln P_{t-1,t} = \ln P_{t,t} - \ln P_{t-1,t}$.
- Rolling volatility: e.g., σ_t from a causal rolling window (std of rr).

- **Triple-barrier labeling:** upper barrier $+\kappa\sigma+k\sigma$, lower barrier $-\kappa\sigma-k\sigma$, and max horizon HH minutes; label = which barrier hits first (or 0 if neither).
 - **Purge:** drop training samples overlapping in time with validation label horizons.
 - **Embargo:** hold-out gap after each validation fold to prevent leakage via close-by outcomes.
 - **Uniqueness / concurrency weights:** down-weight overlapping events.
 - **EV:** expected value per trade net of costs.
-

3) System Architecture (High-Level)

1. **Ingestion & Integrity**
 2. **Bar Construction** (time bars + info-driven bars)
 3. **Feature Engineering** (strictly causal)
 4. **Labeling** (triple-barrier + sample weights)
 5. **Split & Validation** (purged, embargoed, nested walk-forward)
 6. **Model Training** (regularized baselines)
 7. **Probability Calibration**
 8. **Ensembling** (across models & horizons)
 9. **Signal Gating & Execution Mapping** (EV thresholds, hysteresis, cool-downs)
 10. **Risk Controls & Sizing** (vol target, limits, kill-switches)
 11. **Backtesting Engine** (after-cost, event-driven)
 12. **Deployment** (real-time inference)
 13. **Monitoring & Retraining** (drift, calibration, stability)
-

4) Data Requirements

4.1 Inputs (must-have per bar)

- Timestamp (UTC), Open, High, Low, Close, Volume.
- Trades count (if available) and notional volume (quote & base).
- **Fees** (maker/taker schedule), **estimated spread** series (from best bid/ask if available; else proxy).
- Optional: microstructure (order book top-of-book), but not required for v1.

4.2 Timezone & Ordering

- Store raw in UTC. Convert for displays as needed. Bar close at **end of minute**. Ensure strict monotonic timestamps; fill missing minutes explicitly with NaNs then forward engineering **must not** peek into future.

4.3 Data Integrity Checks (blocking)

- No duplicate timestamps.
 - No negative prices/volumes.
 - Gaps identified and logged.
 - Roll-forward/backfill **disallowed** for features—use only causal windows.
-

5) Bar Construction

5.1 Time Bars

- Primary: 1m bars (original).
- Aggregates: 5m, 15m, 60m computed **causally** (downsample by grouping past minutes).

5.2 Information-Driven Bars (optional v1, required v1.1)

- **Volume bars**: emit a bar each time cumulative base volume $\geq \text{threshold } V * V^*$.
 - **Dollar bars**: emit a bar each time cumulative quote notional $\geq \text{threshold } D * D^*$.
 - Thresholds selected to roughly match the typical data density of 1–5m bars.
 - Maintain separate datasets per bar type/horizon; **do not merge** into a single interleaved sequence. Models ensemble across them later.
-

6) Feature Engineering (Causal Only)

6.1 Core Features (Minimal, Robust)

- Returns: $rt, 1r_{\{t,1\}}$ (1m), $rt, 5r_{\{t,5\}}$, $rt, 15r_{\{t,15\}}$ (compounded).
- Volatility: rolling std of 1m returns over windows $W \in \{15, 60, 240\}$.
- Ranges: HL%, CO%, true range proxies (causal).
- Volume/trade intensity: rolling z-scores (causal).
- Skew/kurtosis: rolling higher moments on returns (small windows; clamp outliers).
- **Fractional differencing**: price series fractionally differenced with $d \in [0.2, 0.6]$ tuned to first achieve stationarity on **training only**. Include as features; **do not replace** returns.

6.2 Denoising

- **Allowed**: causal EMA/IIR smoothing on features; **no zero-phase** filters (e.g., `filtfilt`) and **no future look-ahead** (wavelets only if implemented strictly causal; otherwise skip).
- If using Butterworth/IIR: apply single-pass causal form; document order & cutoff; verify no phase compensation from the future.

6.3 Scaling & Leakage Control

- Feature scaling (if used) must be **fit on training folds only** and applied to validation/test.
- Rolling statistics for each t must use data $\leq t$.
- Drop any feature that requires future bars.

6.4 Redundancy Pruning

- Remove features with $|\rho| > 0.95$ within training folds.
- Stability selection: keep features with consistent importance/sign across folds.

7) Labeling — Triple-Barrier Events

7.1 Parameters

- $H \in \{30, 60, 90\}$ minutes (grid).
- $k \in \{1.0, 1.5, 2.0\}$ multiple of rolling σ_t (choose σ window from $\{60, 240\}$ mins).
- Profit barrier: $+k \cdot \sigma_t$; stop barrier: $-k \cdot \sigma_t$.
- Starting price: close at t (or mid if available; be consistent across pipeline).

7.2 Rules

- Simulate forward from t up to $t + Ht + H$ (causal only) to detect first touch of either barrier.
- Label $y_t \in \{+1, -1, 0\}$: +1 if upper hits first; -1 if lower; 0 if neither.
- **Overlap accounting**: compute **event concurrency** (how many active events overlap at each time) and derive **uniqueness weights** $w_t \propto 1/\text{concurrency}_t$; normalize to mean 1 on the training set.

7.3 Class Imbalance

- Use class weights or focal loss; **no oversampling**. Preserve label base rates.

8) Validation Protocol — Purged, Embargoed, Nested Walk-Forward

8.1 Fold Construction

- Use $K=5$ chronological folds. For each fold:
 - **Train** on past window.
 - **Validate** on next segment.
 - **Purge**: remove from training any samples whose event horizon overlaps validation.
 - **Embargo**: skip an additional gap of length H after validation start.

8.2 Nested Tuning

- Inner loop: choose hyperparameters by maximizing **after-cost Sharpe** on the inner validation (same purge/embargo).

- Outer loop: report performance distribution across folds (Sharpe, Sortino, turnover, max DD).
 - **Final test:** the most recent untouched window, reserved from the start; used for paper-trading only. No tuning based on this set.
-

9) Models & Hyperparameters

9.1 Baseline 1 — Logistic Regression (Elastic-Net)

- Inputs: selected features.
- Solver: liblinear/saga.
- Grid: $C \in \{0.01, 0.1, 1, 10\}$; $l1_ratio \in \{0, 0.25, 0.5, 0.75, 1\}$
- Class weights: “balanced” or inverse frequency.
- Use **sample weights** `wtw_t` from uniqueness.

9.2 Baseline 2 — LightGBM

- Objective: binary or multiclass (for $\{-1, 0, +1\}$); recommend **binary one-vs-rest** with meta-labeling later; for v1, binary “act vs no-act” is acceptable if chosen.
- Key constraints (regularization first):
 - `max_depth` $\in \{3, 4, 5\}$
 - `num_leaves` consistent with depth ($\leq 2^{(\text{depth}+1)}$)
 - `min_data_in_leaf` $\in \{200, 500, 1000\}$ (tune vs dataset size)
 - `feature_fraction` $\in \{0.6, 0.8\}$
 - `bagging_fraction` $\in \{0.6, 0.8\}$, `bagging_freq` > 0
 - `lambda_l1` $\in \{0, 1e-3, 1e-2, 1e-1\}$
 - `lambda_l2` $\in \{0, 1e-3, 1e-2, 1e-1\}$
 - `learning_rate` $\in \{0.02, 0.05, 0.1\}$; `early_stopping_rounds`=200; `num_boost_round` up to 5000.
- Use **sample weights** `wtw_t`.
- Monotonic constraints: only if domain knowledge supports (optional).

9.3 Baseline 3 — CatBoost

- Ordered boosting **on**.
- `Depth` $\in \{3, 4, 5\}$; `learning_rate` $\in \{0.02, 0.05, 0.1\}$; `l2_leaf_reg` $\in \{1, 3, 10, 30\}$; `bagging_temperature` $\in \{0, 1\}$; `subsample` $\in \{0.6, 0.8\}$ if not using ordered bootstrap.
- Early stopping: 200 rounds.
- Sample weights: `wtw_t`.

Note: Deep models (TCN/LSTM/Transformers) are **out of scope for v1**; revisit only if v1 is stable and survives robustness tests.

10) Probability Calibration

- Method: **Platt scaling** (logistic) or **isotonic regression**.
 - Fit calibrator **per fold** on the **validation** predictions only; apply to corresponding test segment.
 - Evaluate **Brier score** and reliability curves.
 - Store calibrated ptp_t for execution mapping.
-

11) Ensembling

11.1 Across Models

- Weighted average of calibrated probabilities from **Logit, LightGBM, CatBoost**.
- Weights learned **only on training** via a constrained optimizer (non-negative, sum to 1) to maximize **after-cost Sharpe** on validation; no access to test.
- Simpler fallback: equal weights.

11.2 Across Horizons / Bar Types

- Maintain separate model streams for 1m, 5m, 15m time bars and (when available) dollar/volume bars.
 - Ensemble by averaging probabilities with horizon-specific weights (constrained as above).
 - Optional rule: **regime gating** (e.g., high vol → favor shorter horizons).
-

12) Signal Gating & Execution Mapping

12.1 EV Calculation

For a long signal with $TP = +\theta$ (in bps), $SL = -\theta$, calibrated probability pp of profit barrier hit within HH , and round-trip cost CC (bps):

$$EV = p \cdot \theta - (1-p) \cdot \theta - C = (2p-1)\theta - C \quad \text{CEV} = p \cdot \theta - (1-p) \cdot \theta - C = (2p-1)\theta - C$$

- Enter only if **EV > 0 and EV > M** (safety margin MM , e.g., $2 \times$ median cost over last week).

12.2 Thresholding, Hysteresis, Cool-Downs

- Two thresholds $T_{enter} > T_{exit}$ on pp or EV to reduce churn (hysteresis).
- Require persistence: signal must satisfy threshold for **N consecutive bars** ($N \in \{2,3\}$).
- After exit: **cool-down** of Cd_{d} minutes (e.g., 10–30) before re-entry in same direction.

12.3 Position Sizing & Volatility Target

- Raw size $s^* = \text{clip}(2p-1, 0, 1)s^* = \text{clip}(2p-1, 0, 1)$.
- Scale to meet **volatility target**: choose leverage so that rolling 1-day realized vol of PnL \approx target v_{nu} (e.g., 10% annualized).
- Hard caps: max leverage L_{max} ; max position Q_{max} ; max daily turnover U_{max} .

12.4 Risk Controls & Kill-Switches

- Per-trade stopout at SL; per-day stop after max realized drawdown D_{day} .
- Rolling **max drawdown** limit over last 90 days triggers de-risk (halve size) and at next threshold triggers **all-stop** until manual review.
- Capacity checks: reject trades if expected slippage $>$ threshold.

13) Costs & Slippage Model (Backtest & Live)

- **Fees**: maker/taker per exchange schedule; encode explicitly.
- **Spread**: use best quotes if available; else rolling proxy from high-low micro-ranges.
- **Slippage**: function of volatility \times size; baseline: $\text{slip} = \alpha \cdot \sigma_{1m} \cdot \text{size} / \text{ADV}$. Tune α conservatively.
- **Partial fills**: simulate with queue priority = taker unless resting limit order logic is implemented (not in v1).
- Apply costs at **execution time** (signal \rightarrow order \rightarrow fill) in backtest.

14) Backtesting Engine (Deterministic, Event-Driven)

- **Inputs**: bars, signals, costs, sizing rules.
- **Latency assumption**: orders placed at bar close, filled at **next bar open** price \pm half-spread \pm slippage (configurable).
- **Order types**: market only in v1; limit logic out of scope.
- **Position accounting**: FIFO; PnL includes fees and slippage.
- **Outputs**: per-trade log, per-bar PnL, equity curve.

15) Evaluation Metrics (After Costs Only)

- Annualized **Sharpe** and **Sortino**.
- **Hit rate, average win/loss, expectancy**.
- **Turnover** (daily, annualized).
- **Max drawdown, Calmar/MAR**.
- **Stability**: std of fold Sharpe; inter-quartile range across folds.

- **Calibration:** Brier score, reliability curves.
- **Capacity:** slippage vs size sensitivity.

Model selection optimizes for **after-cost Sharpe** subject to constraints: turnover \leq target, max DD \leq cap, calibration error \leq threshold.

16) Robustness & Stress Testing

- **Block/bootstrap:** stationary/block bootstrap of returns to test persistence.
- **Cost stress:** $\pm(2-3) \times$ fees/spread/slippage.
- **Parameter knockouts:** perturb key HPs $\pm 20\%$ to test brittleness.
- **Regime slicing:** high vs low volatility, trend vs range periods.
- **p-value style checks:** reality-check style multiple testing control (at least report selection count and adjusted expectations).

Pass/fail criteria: strategy maintains positive after-cost Sharpe and acceptable drawdowns in at least **80%** of stress scenarios.

17) Deployment Requirements

17.1 Real-Time Inference

- Inference cycle: every minute at bar close + 1–2s buffer.
- Max end-to-end latency (feature calc \rightarrow signal \rightarrow order): **< 3 seconds**.
- Deterministic feature state managed via a simple **feature store** (rolling windows updated incrementally).

17.2 Packaging & Reproducibility

- Fixed random seeds per fold and model.
- Config-driven (YAML): data paths, HP grids, thresholds, costs.
- Version everything: code, configs, datasets (hash/snapshot).
- Experiment tracker (any): log HPs, metrics, artifacts (models, calibrators).

17.3 Monitoring

- Live metrics: rolling Sharpe, turnover, hit rate, slippage vs model, calibration drift (Brier), latency.
- Alerts: breach of kill-switch thresholds, missing data, latency > SLA, calibration drift > tolerance.
- Logging: per-decision record (features hash, model version, probs, EV, action, fill).

17.4 Retraining

- Cadence: weekly (or after **N** new days or regime change trigger).

- **Process:** freeze last live model; retrain via the exact pipeline; validate; paper-trade 1–2 weeks before promotion.
 - **Rollback:** ability to revert to last stable model instantly.
-

18) Implementation Checklist (Top-Down)

1. **Data ingestion & integrity checks** done.
 2. **Bar builders** (1m, 5m, 15m; + volume/dollar in v1.1) ready.
 3. **Causal features** implemented; leakage tests passed.
 4. **Triple-barrier labels + uniqueness weights** implemented.
 5. **Purged+embargoed nested walk-forward** splitter implemented and unit-tested.
 6. **Baselines** (Logit, LGBM, CatBoost) train with **sample weights** and early stopping.
 7. **Calibration** per fold (Platt/Isotonic) with stored calibrators.
 8. **Ensemble** weights fitted on training folds only; applied out-of-fold.
 9. **Execution mapping** (EV thresholds, hysteresis, cool-downs) implemented.
 10. **Vol target sizing, risk caps, kill-switches** wired.
 11. **Backtester** simulates orders with costs and outputs trade logs + metrics.
 12. **Reports:** cross-fold summaries, robustness suite, final untouched paper-trade report.
 13. **Deployment** scaffold: inference loop, monitoring, alerting.
 14. **Runbook** for retraining & rollback.
-

19) Acceptance Criteria

- **Leakage tests:** unit tests that deliberately try to use future info must fail; all rolling ops verified causal.
 - **Validation protocol:** code proves purge/embargo lengths equal to label horizon HH.
 - **Metrics:** on cross-fold OOS, after-cost Sharpe > 0 with stable dispersion; calibration Brier < naive baseline; turnover within configured limit.
 - **Stress suite:** positive after-cost Sharpe in ≥ 80% scenarios; drawdowns under cap.
 - **Paper-trade window** (most recent period): after-cost Sharpe within **50–100%** of cross-fold median; turnover and DD consistent.
 - **Monitoring:** live dashboards and alerts configured; dry-run shows signals/alerts firing as expected.
 - **Reproducibility:** re-running with same seeds reproduces metrics within statistical jitter.
-

20) Micro-Goals by Phase (Why each exists)

- **Labels & Splits:** Remove noise and leakage → trustworthy edges.
 - **Features:** Small, causal, stationarized set → lower variance.
 - **Models:** Regularized baselines → strong generalization.
 - **Calibration:** Turn scores into **actionable probabilities** for EV mapping.
 - **Ensembles:** Reduce variance across models/horizons.
 - **Execution & Risk:** Convert edge to PnL while controlling churn and drawdowns.
 - **Backtest Realism:** Prevent cost/latency fantasies.
 - **Monitoring & Retraining:** Catch drift and keep the edge alive.
-

21) Guardrails & Non-Negotiables

- **Never** use non-causal filters (filtfilt, centered windows, future-aware transforms).
 - **Never** tune on the final test/paper-trade window.
 - **Always** evaluate **after costs**; report costs separately.
 - **Document** every assumption (latency, slippage, fees, horizons).
-

22) Deliverables

1. **Code package** with modules: ingestion, bars, features, labels, cv_splitter, models, calibration, ensemble, execution, risk, backtester, reports, deploy.
 2. **Config files** (YAML) with all parameters and grids.
 3. **Unit tests** (features causality, purge/embargo, calibration separation).
 4. **Integration tests** (end-to-end backtest on a small dataset).
 5. **Experiment logs** and **artifacts** (models + calibrators per fold).
 6. **Backtest report** (cross-fold + stress + paper-trade).
 7. **Runbooks** (deployment, monitoring, retraining, rollback).
-

23) Optional Extensions (Post-v1)

- **Meta-labeling:** a secondary classifier deciding when to act on primary signals.
- **Uncertainty gating** via **conformal prediction** or CV dispersion thresholds.
- **Regime models:** state detection to select horizon weights.
- **Info-driven bars** (volume/dollar) moved from optional to default after evaluation.

- **Deep models** (TCN) under same validation strictness; keep only if they beat baselines **after costs** and survive robustness tests.
- **RL**: only with a high-fidelity simulator and strict OOS controls.

24) Appendix — Pseudocode Outlines

24.1 Triple-Barrier Labeling

for each index t:

start_price = Close[t]

upper = start_price * exp(+k * sigma_t)

lower = start_price * exp(-k * sigma_t)

for tau in 1..H:

p = Close[t+tau] # or mid, strictly future-only within label calc

if p >= upper: y[t] = +1; T_end[t] = t+tau; break

if p <= lower: y[t] = -1; T_end[t] = t+tau; break

if no barrier hit by H: y[t] = 0; T_end[t] = t+H

Compute concurrency over all active [t, T_end[t]]; uniqueness weight wtw_t inversely proportional to concurrency, normalized.

24.2 Purged, Embargoed Walk-Forward

split timeline into K folds chronologically

for fold i:

val = segment_i

embargo_gap = H

train = data before val, excluding:

- any t whose [t, T_end[t]] overlaps val
- embargo_gap minutes after val start

fit inner-CV for HPs on train

evaluate on val

24.3 EV Thresholding & Sizing

p = calibrated_prob_long

theta = k * sigma_t_in_bps

EV = (2*p - 1) * theta - Cost

if EV > max(0, M) and persistence_condition_met and not in_cooldown:

```
size_raw = clip(2*p - 1, 0, 1)

size = vol_target(size_raw)

place_order(size)

else:

    hold_or_exit()
```

25) Why We Integrated Each Approach

- **Triple-barrier labeling:** focuses learning on **meaningful moves**, reducing label noise from 1-min chops.
 - **Purged & embargoed CV:** eliminates leakage from overlapping horizons; realistic OOS estimates.
 - **Downsampling & info-driven bars:** improve signal-to-noise; we **add** them as parallel views to ensemble, not as a single replacement.
 - **Causal smoothing & fractional differencing:** stabilize features while preserving memory and avoiding leakage.
 - **Regularized baselines (Logit/LGBM/CatBoost):** strong on noisy tabular data with built-in variance control.
 - **Probability calibration:** converts raw scores into **trustworthy probabilities** for EV mapping and sizing.
 - **Ensembling across models/horizons:** reduces variance and regime sensitivity.
 - **EV gating, hysteresis, cool-downs, vol-targeting:** translates statistical edge into **tradable** performance with controlled churn and risk.
 - **Robustness & stress tests:** filter out backtest mirages and parameter accidents before capital is at risk.
-

End of spec.

If you need this converted into a project scaffold (folder structure, config templates, and test stubs), say the word and I'll lay it out.